

Testing Monotonicity

Oded Goldreich*

Shafi Goldwasser†

Eric Lehman†

Dana Ron‡

Abstract

We present a (randomized) test for monotonicity of Boolean functions. Namely, given the ability to query an unknown function $f : \{0, 1\}^n \mapsto \{0, 1\}$ at arguments of its choice, the test always accepts a monotone f , and rejects f with high probability if it is ϵ -far from being monotone (i.e., every monotone function differs from f on more than an ϵ fraction of the domain). The complexity of the test is $\text{poly}(n/\epsilon)$.

The analysis of our algorithm relates two natural combinatorial quantities that can be measured with respect to a Boolean function; one being global to the function and the other being local to it.

We also consider the problem of testing monotonicity based only on random examples labeled by the function. We show an $\Omega(\sqrt{2^n/\epsilon})$ lower bound on the number of required examples, and provide a matching upper bound (via an algorithm).

1 Introduction

In this work we address the problem of *testing whether a given Boolean function is monotone*. A function $f : \{0, 1\}^n \mapsto \{0, 1\}$ is said to be monotone if $f(x) \leq f(y)$ for every $x \prec y$, where \prec denotes the natural partial order among strings (i.e., $x_1 \cdots x_n \prec y_1 \cdots y_n$ if $x_i \leq y_i$ for every i and $x_i < y_i$ for some i). The testing algorithm can request the value of the function on arguments of its choice, and is required to distinguish monotone functions from functions that are far from being monotone.

More precisely, the testing algorithm is given a *distance* parameter $\epsilon > 0$, and oracle access to an unknown function f mapping $\{0, 1\}^n$ to $\{0, 1\}$. If f is a monotone then the algorithm should accept it with probability at least $2/3$, and if f is at distance greater than ϵ from any monotone function then the algorithm should reject it with probability at least $2/3$. Distance between functions is measured in terms of the fraction of the domain on which the functions differ. The

complexity measures we focus on are the *query complexity* and the *running time* of the testing algorithm.

We present a randomized algorithm for testing the monotonicity property whose query complexity and running time are polynomial in n and $1/\epsilon$. The algorithm performs a simple local test: It verifies whether monotonicity is maintained for randomly chosen pairs of strings that differ exactly on a single bit. In our analysis we relate this local measure to the global measure we are interested in — the minimum distance of the function to any monotone function.

1.1 Perspective

Property Testing, as explicitly defined by Rubinfeld and Sudan [28] and extended in [19], is best known by the special case of *low degree testings* [12, 17, 28, 27, 4] which plays a central role in the construction of probabilistically checkable proofs (PCP) [6, 5, 16, 3, 2, 27, 4]. The recognition that property testing is a general notion has been implicit in the context of PCP: It is understood that low degree tests as used in this context are actually codeword tests (in this case of BCH codes), and that such tests can be defined and performed also for other error-correcting codes such as the Hadamard code [2, 9, 10, 7, 8, 26, 29], and the “Long Code” [8, 22, 23, 29].

Forasmuch as error-correcting codes emerge naturally in the context of PCP, they do not seem to provide a natural representation of familiar objects whose properties we may wish to investigate. That is, one can certainly encode any given object by an error-correcting code — resulting in a (legitimate yet) probably unnatural representation of the object — and then test properties of the encoded object. However, this can hardly be considered as a “natural test” of a “natural phenomena”. For example, one may indeed represent a graph by applying an error correcting code to its adjacency matrix (or to its incidence list), but the resulting string is not the “natural representation” of the graph.

The study of Property Testing as applied to natural representation of (non-algebraic) objects was initiated in [19]. In particular, Property Testing as applied to *graphs* has been studied in [19, 20, 21] — where the first work considers the *adjacency matrix representation* of graphs (most adequate for dense graphs), and the latter works consider the *incidence list representation* (adequate for sparse graphs).

In this work we consider property testing as applied to the most generic (i.e., least structured) object — an arbitrary

*Department of Computer Science and Applied Mathematics, Weizmann Institute of Science, Rehovot, ISRAEL.
E-mail: oded@wisdom.weizmann.ac.il. Work done while visiting LCS, MIT.

†Laboratory for Computer Science, MIT, 545 Technology Sq., Cambridge, MA 02139. E-mail: {shafi,e.lehman,danar}@theory.lcs.mit.edu.

‡Supported by an ONR Science Scholar Fellowship at the Bunting Institute.

Boolean function. In this case the choice of representation is “forced” upon us.

1.2 Monotonicity

In interpreting monotonicity it is useful to view Boolean functions over $\{0, 1\}^n$ as subsets of $\{0, 1\}^n$, called *concepts*. This view is the one usually taken in the PAC Learning literature. Each position in $\{1, \dots, n\}$ corresponds to a certain *attribute*, and a string $x = x_1 \dots x_n \in \{0, 1\}^n$ represents an instance where $x_i = 1$ if and only if the instance x has the i^{th} attribute. Thus, a concept (subset of instances) is monotone if the presence of additional attributes maintains membership of instances in the concept (i.e., if instance x is in the concept C then any instance resulting from x by adding some attributes is also in C).

The class of monotone concepts is quite general and rich. On the other hand, monotonicity suggests a certain aspect of simplicity. Namely, each attribute has a uni-directional effect on the value of the function. Thus, knowing that a concept is monotone may be useful in various applications. In fact, this form of simplicity is exploited by Angluin’s learning algorithm for monotone concepts [1], which uses membership queries and has complexity that is linear in the number of terms of the target concept’s DNF representation.

We note that an efficient tester for monotonicity is useful as a preliminary stage before employing Angluin’s algorithm. As is usually the case, Angluin’s algorithm relies on the premise that the unknown target concept is in fact monotone. It is possible to simply apply the learning algorithm without knowing whether the premise holds, and hope that either the algorithm will succeed nonetheless in finding a good hypothesis or detect that the target is not monotone. However, due to the dependence of the complexity of Angluin’s algorithm on the number of terms of the target concept’s DNF representation, it may be much more efficient to first test whether the function is at all monotone (or close to it).

1.3 The natural monotonicity test

The main result of the paper is that a tester for monotonicity is obtained by repeating the following for $\text{poly}(n/\epsilon)$ many times: Uniformly select a pair of strings at Hamming distance 1 and check if monotonicity is satisfied with respect to the value of f on these two strings. That is,

ALGORITHM 1: On input n, ϵ and oracle access to $f : \{0, 1\}^n \mapsto \{0, 1\}$, repeat the following steps up to n^3/ϵ times

1. Uniformly select $x \in \{0, 1\}^n$ and $i \in \{1, \dots, n\}$.
2. Obtain the values of $f(x)$ and $f(y)$, where y results from x by flipping the i^{th} bit.
3. If $x, y, f(x), f(y)$ demonstrate that f is not monotone then reject.

That is, if either $(x \prec y) \wedge (f(x) > f(y))$ or $(y \prec x) \wedge (f(y) > f(x))$ then reject.

If all iterations were completed without rejecting then accept.

Theorem 1 (main result): *Algorithm 1 is a testing algorithm for monotonicity. Furthermore, if the function is monotone then Algorithm 1 always accepts.*

Theorem 1 asserts that a (random) *local check* (i.e., Step 3 above) can establish the existence of a *global property* (i.e., the distance of f to the set of monotone functions). Actually, Theorem 1 is proven by relating two quantities referring to the above: Given $f : \{0, 1\}^n \mapsto \{0, 1\}$, we denote by $\delta_M(f)$ the fraction of pairs (x, y) in which Step 3 rejects. Observe that $\delta_M(f)$ is actually a combinatorial quantity (i.e., the fraction of pairs of n -bit strings, differing on one bit, which violate the monotonicity condition). We then define $\epsilon_M(f)$ to be the distance of f from the set of monotone functions (i.e., the minimum over all monotone functions g of $|\{x : f(x) \neq g(x)\}|/2^n$). Observing that Algorithm 1 always accepts a monotone function, Theorem 1 follows from Theorem 2, stated below.

Theorem 2 *For any $f : \{0, 1\}^n \mapsto \{0, 1\}$,*

$$\delta_M(f) \geq \frac{\epsilon_M(f)}{n^3}.$$

We comment that a slightly more careful analysis yields a better bound than the one stated in the theorem: namely,

$$\delta_M(f) = \Omega\left(\frac{\epsilon_M(f)}{n^2 \log(1/\epsilon_M(f))}\right). \quad (1)$$

As for the reverse direction; that is, lower bounding $\epsilon_M(f)$ in terms of $\delta_M(f)$, we have

Proposition 3 *For every function $f : \{0, 1\}^n \mapsto \{0, 1\}$, $\epsilon_M(f) \geq \delta_M(f)/2$.*

Thus, for every function f

$$\frac{\epsilon_M(f)}{\text{poly}(n)} \leq \delta_M(f) \leq O(\epsilon_M(f))$$

A natural question that arises is that of the exact relation between $\delta_M(\cdot)$ and $\epsilon_M(\cdot)$. We observe that this relation is not simple; that is, it does not depend only on the values of δ_M and ϵ_M .

Proposition 4 *The following holds for every n and every $2^{-c \cdot n} \leq \alpha \leq \frac{1}{2} - O(\frac{1}{\sqrt{n}})$, where c is any constant strictly smaller than 1.*

1. *There exists a function $f : \{0, 1\}^n \mapsto \{0, 1\}$ such that $\alpha \leq \epsilon_M(f) \leq 2\alpha$ and $\delta_M(f) = \Theta\left(\frac{\epsilon_M(f)}{\sqrt{n}}\right)$.*

2. *There exists a function $f : \{0, 1\}^n \mapsto \{0, 1\}$ such that $\alpha \leq \epsilon_M(f) \leq 2\alpha$ and $\delta_M(f) = \Theta(\epsilon_M(f))$.*
3. *For any $\alpha = O(n^{-\frac{3}{2}})$, there exists a function $f : \{0, 1\}^n \mapsto \{0, 1\}$ such that $\alpha \leq \epsilon_M(f) \leq 2\alpha$ and $\delta_M(f) = \Theta\left(\frac{\epsilon_M(f)}{n}\right)$.*

PERSPECTIVE. Analogous quantities capturing local and global properties of functions were analyzed in the context of *linearity testing*. For a function $f : \{0, 1\}^n \mapsto \{0, 1\}$ (as above), one may define $\epsilon_{\text{LIN}}(f)$ to be its distance from the set of linear functions and $\delta_{\text{LIN}}(f)$ to be the fraction of pairs, $(x, y) \in \{0, 1\}^n \times \{0, 1\}^n$ for which $f(x) + f(y) \neq f(x \oplus y)$. A sequence of works [12, 9, 10, 7] has demonstrated a fairly complex behavior of the relation between δ_{LIN} and ϵ_{LIN} . The interested reader is referred to [7].

1.4 Testing based on random examples

Algorithm 1 makes essential use of queries. We show that this is no coincidence – any monotonicity tester that utilizes only uniformly and *independently* chosen random examples, must have much higher complexity.

Theorem 5 *For any $\epsilon = O(n^{-3/2})$, any tester for monotonicity that only utilizes random examples must use at least $\Omega(\sqrt{2^n/\epsilon})$ such examples.*

Interestingly, this lower bound is tight up to a $\text{poly}(n)$ factor.

Theorem 6 *There exists a tester for monotonicity which only utilizes random examples and uses at most $O(\sqrt{n^3 \cdot 2^n/\epsilon})$ examples. Furthermore, the algorithm runs in time $\text{poly}(n) \cdot \sqrt{2^n/\epsilon}$.*

We note that the above tester is significantly faster than any learning algorithm for the class of all monotone concepts when the allowed error is $O(1/\sqrt{n})$: Learning (under the uniform distribution) requires $\Omega(2^n/\sqrt{n})$ examples (and even that number of queries) [24].¹

1.5 Extensions and Open Problems

TESTING UNATENESS. A function $f : \{0, 1\}^n \mapsto \{0, 1\}$ is said to be unate if for every x_i (where $x = x_1 \dots x_n$ is the input to the function), exactly one of the following holds: whenever the value of x_i is flipped from 0 to 1 then the value of f does not

decrease; or whenever the value of x_i is flipped from 1 to 0 then the value of f does not decrease. Thus, unateness is a more general notion than monotonicity. We show that our algorithm can be extended to test whether a Boolean function is unate or far from any unate function. The query and time complexities of the (extended) algorithm are bounded by $O(n^{3.5}/\epsilon)$.

OTHER DOMAINS AND RANGES. Let Σ and Ξ be finite sets, and $<_\Sigma$ and $<_\Xi$ (total) orders on Σ and Ξ , respectively. Then we can extend the notion of monotonicity to functions from Σ^n to Ξ , in the obvious manner: Namely, a function $f : \Sigma^n \mapsto \Xi$ is said to be monotone if $f(x) \leq_\Xi f(y)$ for every $x \prec_\Sigma y$, where $x_1 \dots x_n \prec_\Sigma y_1 \dots y_n$ if $x_i \leq_\Sigma y_i$ for every i and $x_i <_\Sigma y_i$ for some i . Our algorithm generalizes to testing monotonicity over extended domains and ranges. The complexity of the generalized algorithm scales quadratically with $|\Sigma|$ and linearly with $|\Xi|$. It is an interesting open problem whether these dependencies can be removed (or reduced). In particular, we believe that the dependence on the size of the range Ξ can be removed (for more details, see discussion in the full version of this paper [18]).

REMOVING THE DEPENDENCE ON n . Our algorithm (even for the base case), has a polynomial dependence on the dimension of the input, n . As shown in Proposition 4, some dependence of the query complexity on n is unavoidable in the case of our algorithm. However, it is an interesting open problem whether other algorithms may have significantly lower query (and time) complexities, and in particular have query complexity independent of n . A candidate alternative algorithm inspects pairs of strings x, y , where x is chosen uniformly in $\{0, 1\}^n$, and y is chosen as follows: First select an index (weight) $w \in \{0, \dots, n\}$ with probability $\binom{n}{w} \cdot 2^{-n}$, and then select y uniformly among the strings having w 1's, and being comparable to x (i.e., $y \prec x$ or $y \succ x$).

Related Work

The “spot-checker for sorting” presented in [14, Sec. 2.1] implies a tester for monotonicity with respect to functions from any fully ordered domain to any fully ordered range, having query and time complexities that are logarithmic in the size of the domain. We note that this problem corresponds to the special case of $n = 1$ of the extension discussed in Subsection 1.5 (to general domains and ranges).

Organization

Theorem 2 is proved in Section 3. The extension to testing unateness is presented in Section 4. Propositions 3 and 4 and Theorems 5 and 6, as well as the other extensions, are provided in the full version of this paper [18].

¹ The claim follows by considering all possible concepts that contain all instances having $\lfloor n/2 \rfloor + 1$ or more 1's, no instances having $\lfloor n/2 \rfloor - 1$ or less 1's, and any subset of the instances having exactly $\lfloor n/2 \rfloor$ 1's. In contrast, “weak learning” [25] is possible in polynomial time. Specifically, the class of monotone concepts can be learned in polynomial time with error at most $1/2 - \Omega(1/\sqrt{n})$ (though no polynomial-time learning algorithm can achieve an error of $1/2 - \omega(\log(n)/\sqrt{n})$) [11].

2 Preliminaries

For any pair of functions $f, g : \{0, 1\}^n \rightarrow \{0, 1\}$, we define the *distance* between f and g , denoted, $\text{dist}(f, g)$, to be the fraction of instances $x \in \{0, 1\}^n$ on which $f(x) \neq g(x)$. In other words, $\text{dist}(f, g)$ is the probability over a uniformly chosen x that f and g differ on x . Thus, $\epsilon_M(f)$ as defined in the introduction is the minimum, taken over all monotone functions g of $\text{dist}(f, g)$.

A general formulation of Property Testing was suggested in [19], but here we consider a special case formulated previously in [28].

Definition 1 (property tester): *Let $P = \cup_{n \geq 1} P_n$ be a subset (or a property) of Boolean functions, so that P_n is a subset of the functions mapping $\{0, 1\}^n$ to $\{0, 1\}$. A (property) tester for P is a probabilistic oracle machine², M , which given n , a distance parameter $\epsilon > 0$ and oracle access to an arbitrary function $f : \{0, 1\}^n \rightarrow \{0, 1\}$ satisfies the following two conditions:*

1. The tester accepts f if it is in P :

If $f \in P_n$ then $\text{Prob}(M^f(n, \epsilon) = 1) \geq \frac{2}{3}$.

2. The tester rejects f if it is far from P :

If $\text{dist}(f, g) > \epsilon$ for every $g \in P_n$, then $\text{Prob}(M^f(n, \epsilon) = 1) < \frac{1}{3}$.

TESTING BASED ON RANDOM EXAMPLES. In case the queries made by the tester are uniformly and *independently* distributed in $\{0, 1\}^n$, we say that it only uses examples. Indeed, a more appealing way of looking as such a tester is as an ordinary algorithm (rather than an oracle machine) which is given as input a sequence $(x_1, f(x_1)), (x_2, f(x_2)), \dots$ where the x_i 's are uniformly and independently distributed in $\{0, 1\}^n$.

Definition 2 (the Boolean-Lattice graph): *For every string $x \in \{0, 1\}^n$, let $w(x)$ denote the weight of x (i.e., the number of 1's in x). For each i , $0 \leq i \leq n$, let $L_i \subset \{0, 1\}^n$ denote the set of n -bit strings of weight i (i.e., $L_i = \{x \in \{0, 1\}^n : w(x) = i\}$). Let G_n be the leveled directed (acyclic) graph over the vertex set $\{0, 1\}^n$, where there is a directed edge from y to x if and only if $x \prec y$ and $w(x) = w(y) - 1$ (i.e., x and y are in adjacent L_i 's).*

Given the definition of G_n we may view our algorithm as uniformly selecting *edges* in G_n and querying the function f on their end-points. We call an edge directed from y to x in G_n a *violating edge with respect to f* if $f(x) > f(y)$ (whereas $x \prec y$). Thus, $\delta_M(f)$, as defined in the introduction, is the fraction of violating edges in G_n with respect to f .

² Alternatively, one may consider a RAM model of computation, in which trivial manipulation of domain and range elements (e.g., reading/writing an element and comparing elements) is performed at unit cost.

3 Proof of the Main Technical Result

In order to prove Theorem 2 we prove the following two lemmas. The first lemma shows the existence of a *matching* between two relatively large (with respect to $\epsilon_M(f)$) sets of vertices (strings) belonging to different layers of G_n where each vertex y in the first set is matched to a vertex x such that $x \prec y$ but $f(x) > f(y)$. The second lemma shows that for any such matching there exist vertex disjoint (directed) paths in G_n between the two sets (though the paths may correspond to a different matching).

Lemma 7 (existence of large violating matched sets) *For any function $f : \{0, 1\}^n \rightarrow \{0, 1\}$, there exist two sets of vertices $S \subseteq L_s$ and $R \subseteq L_r$, where $s > r$, for which the following holds:*

1. $|S| = |R| \geq \frac{\epsilon_M(f)}{2n^2} \cdot 2^n$;
2. For every $y \in S$, $f(y) = 0$, and for every $x \in R$, $f(x) = 1$;
3. There exists a one-to-one mapping ϕ from S to R such that for every $y \in S$, $\phi(y) \prec y$.

Lemma 8 (existence of disjoint paths between matched sets)

Let r and s be integers satisfying, $0 \leq r < s \leq n$, and let $S \subseteq L_s$ and $R \subseteq L_r$ be sets each of size m . Suppose that there exists a 1-to-1 mapping ϕ from S to R such that for every $y \in S$, there is a directed path in G_n from y to $\phi(y)$. Then there exist m vertex-disjoint directed paths from S to R in G_n .

We prove the two lemmas in the next two subsections. But first we show that Theorem 2 follows by combining the two lemmas.

Proof of Theorem 2: Fixing f we first invoke Lemma 7 to obtain the two matched sets S and R of size at least $m = \frac{\epsilon_M(f)}{2n^2} \cdot 2^n$. By Lemma 8 this matching implies the existence of m vertex disjoint paths from S to R . Consider any such path $z_0 = y, \dots, z_d = x$, where $y \in S$, $x \in R$, and $d = s - r$. Since $z_0 \in S$, we have $f(z_0) = 0$. On the other hand, since $z_d \in R$, we have $f(z_d) = 1$. Therefore, there must exist some $\ell \in \{0, \dots, d-1\}$, such that $f(z_\ell) = 0$ and $f(z_{\ell+1}) = 1$. Thus the edge directed from z_ℓ to $z_{\ell+1}$ is a violating edge with respect to f . Since the paths from S to R are vertex disjoint, they are necessarily edge disjoint, and hence there are at least $m = \frac{\epsilon_M(f)}{2n^2} \cdot 2^n$ such violating edges (at least one per path). Because each vertex in G_n has total degree (indegree plus outdegree) n , the number of edges in G_n is $\frac{1}{2} \cdot 2^n \cdot n$. Therefore, the fraction of violating edges is at least $\frac{\epsilon_M(f)}{n^3}$, and the theorem follows. ■

The strengthening of Theorem 2 stated in Equation (1) is justified by the fact that one may actually show that there exist sets S and R as in Lemma 7 such that $|S| = |R| = \Omega\left(\frac{\epsilon_M(f)}{n \log(1/\epsilon_M(f))}\right) \cdot 2^n$. We show how this improvement can be obtained after we prove Lemma 7.

3.1 Proving the existence of large violating matched sets

Fixing f , let g be a monotone function (over $\{0, 1\}^n$) for which $\text{dist}(f, g) = \epsilon_M(f)$. Namely, g is a monotone function that is closest to f . For $b \in \{0, 1\}$, let

$$D_b \stackrel{\text{def}}{=} \{x : f(x) \neq g(x) \text{ and } g(x) = b\} \quad (2)$$

That is, the set $D_0 \cup D_1$ is a set of minimum size such that if we flip the value of f on all elements in the set then we obtain a monotone function (i.e., g). Since $|D_0 \cup D_1| = \epsilon_M(f) \cdot 2^n$ and $D_0 \cap D_1 = \emptyset$, we may assume, without loss of generality, that $|D_1| \geq \frac{\epsilon_M(f)}{2} \cdot 2^n$. Recall that, by definition,

$$D_1 = \{x : g(x) = 1 \text{ and } f(x) = 0\} \subseteq \{x : f(x) = 0\}$$

For any set $Y \subseteq \{0, 1\}^n$, the shadow³ of Y , denoted $\sigma(Y)$, is defined as follows:

$$\sigma(Y) \stackrel{\text{def}}{=} \{x \notin Y : \exists y \in Y \text{ s.t. } x \prec y\} \quad (3)$$

Namely, the shadow of Y is the set of all strings not in Y that are each smaller than some string in Y . For any $Y \subseteq D_1$ define

$$\sigma_1(Y) \stackrel{\text{def}}{=} \{x \in \sigma(Y) : f(x) = g(x) = 1\} \quad (4)$$

Namely, $\sigma_1(Y)$ is the subset of the shadow of Y containing all strings on which both f and g have value 1. (Note that for any $Y \subseteq D_1$, $\sigma(Y) \setminus \sigma_1(Y) \subseteq \{x : g(x) = 0\}$.) As a visualization (see Figure 3.1), we view g as defining a *boundary* in the Boolean Lattice (similarly, in G_n), such that all strings on and above the boundary are labeled 1, and all other strings are labeled 0. The set D_1 contains those strings above the boundary that f labels 0. The set $\sigma_1(D_1)$ contains all strings in the shadow of D_1 that lie above the boundary. These strings are labeled 1 by f (as otherwise they would be in D_1).

Thus, by definition of D_1 and $\sigma_1(D_1)$, we have that for every $x \in \sigma_1(D_1)$, there exists $y \in D_1$ such that the pair (x, y) satisfies: $x \prec y$ and $f(y) < f(x)$ (i.e., $f(y) = 0$ and $f(x) = 1$). We next show that a stronger statement holds.

Lemma 9 *For every $Y \subseteq D_1$, there exists a 1-to-1 mapping ϕ from Y into $\sigma_1(Y)$, such that for each $y \in Y$, $\phi(y) \prec y$.*

Lemma 9 is the main step in proving Lemma 7 (which also requires that all elements in the set S belong to the same layer in G_n , and that the same hold for all the elements they are mapped to).

Proof: We first show that for every $Y \subseteq D_1$, $|\sigma_1(Y)| \geq |Y|$. Assume towards contradiction that, for some $Y \subseteq D_1$,

³This is not the standard definition of a shadow, as in [13, Chap. 5].




Figure 1. The sets D_1 and $\sigma_1(D_1)$.

$|\sigma_1(Y)| < |Y|$. We show, contrary to our hypothesis on g , that there exists another monotone function g' that is (strictly) closer to f .

Define g' as follows: For every $x \in Y \cup \sigma(Y)$, $g'(x) = 0$. Otherwise, $g'(x) = g(x)$.

We need to verify the following two claims.

Claim 9.1: g' is a monotone function.

Claim 9.2: $\text{dist}(f, g') < \text{dist}(f, g)$.

Proof of Claim 9.1: We need to show that for every x, y such that $x \prec y$, it holds that $g'(x) \leq g'(y)$. Consider the following cases.

Case 1: $x \in Y \cup \sigma(Y)$. In this case $g'(x) = 0$, and so $g'(x) \leq g'(y)$ for all y .

Case 2: $x \notin Y \cup \sigma(Y)$. Note that in this case $g'(x) = g(x)$. We will show that for every y if $x \prec y$ then $y \notin Y \cup \sigma(Y)$ as well, and thus $g'(y) = g(y) \geq g(x) = g'(x)$ as required. Suppose towards contradiction that for some $y \in Y \cup \sigma(Y)$ it holds that $x \prec y$. We consider two subcases.

1. If $y \in Y$ then since $x \prec y$ we have that $x \in Y \cup \sigma(Y)$ in contradiction to the case hypothesis.
2. If $y \in \sigma(Y)$ then there exists $z \in Y$ such that $y \prec z$. Using $x \prec y$ it follows that $x \prec z$ and so again $x \in Y \cup \sigma(Y)$ in contradiction to the case hypothesis.

Claim 9.1 follows. \square

Proof of Claim 9.2: By definition of g' , the functions g and g' differ on the set of strings $\Delta \stackrel{\text{def}}{=} (Y \cup \sigma(Y)) \cap \{x : g(x) = 1\}$. Since $Y \subseteq D_1 \subseteq \{x : g(x) = 1\}$, we have

$$\Delta = Y \cup (\sigma(Y) \cap \{x : g(x) = 1\})$$

$$\begin{aligned}
&= Y \bigcup (\sigma(Y) \cap \{x : g(x) = 1 \text{ and } f(x) = 1\}) \\
&\quad \bigcup (\sigma(Y) \cap \{x : g(x) = 1 \text{ and } f(x) = 0\}) \\
&= Y \bigcup \sigma_1(Y) \bigcup A
\end{aligned}$$

where $A \stackrel{\text{def}}{=} \sigma(Y) \cap \{x : g(x) = 1 \text{ and } f(x) = 0\}$. Consider the three (disjoint) subsets of Δ : Y , $\sigma_1(Y)$, and A .

- For every $x \in Y$, we have $f(x) = 0$ and $g(x) = 1$ (since $Y \subseteq D_1$), and $g'(x) = 0$ (by definition).
Such x contributes to $\text{dist}(f, g)$ but not to $\text{dist}(f, g')$.
- For every $x \in \sigma_1(Y)$, we have $f(x) = g(x) = 1$ (by definition of $\sigma_1(Y)$), and again $g'(x) = 0$.
Such x do not contribute to $\text{dist}(f, g)$ but do contribute to $\text{dist}(f, g')$.
- For every $x \in A$, we have $f(x) = 0$ and $g(x) = 1$ (by definition of A), and again $g'(x) = 0$.
Such x contribute to $\text{dist}(f, g)$ but not to $\text{dist}(f, g')$.

Thus,

$$\begin{aligned}
2^n \cdot (\text{dist}(f, g') - \text{dist}(f, g)) &= |\sigma_1(Y)| - |Y \cup A| \\
&\leq |\sigma_1(Y)| - |Y| < 0
\end{aligned}$$

where the strict inequality is due to the assumption that $|\sigma_1(Y)| < |Y|$. Claim 9.2 follows. \square

Consider any set $Y \subseteq D_1$. We have established that for every $Y' \subseteq Y$, $|\sigma_1(Y')| \geq |Y'|$. Lemma 9 follows from Hall's Theorem (cf. [15, Thm. 6.12]): Consider the auxiliary bipartite graph B whose vertex set is labeled by the strings in $Y \cup \sigma_1(Y)$, and whose edge set is $\{(x, y) : x \in \sigma_1(Y), y \in Y, x \prec y\}$. By the above, for each $Y' \subseteq Y$, we have $|\Gamma(Y')| \geq |Y'|$, where $\Gamma(Y')$ denotes the neighbor set of Y' in B . By Hall's Theorem, this implies that there exists a perfect matching between Y and a subset of $\sigma_1(Y)$. Lemma 9 follows. \blacksquare

Proof of Lemma 7: As noted previously, we may assume that D_1 (see Eq. (2)) has size at least $\epsilon_M(f) \cdot 2^{n-1}$ (the case $|D_0| \geq \epsilon_M(f) \cdot 2^{n-1}$ is analogous). Let $Y_i \stackrel{\text{def}}{=} D_1 \cap L_i$, and let s denote the index of the largest set among the Y_i 's. It follows that $|Y_s| \geq \frac{\epsilon_M(f)}{2^n} \cdot 2^n$.

We now invoke Lemma 9 with $Y = Y_s$. Let $X_s \stackrel{\text{def}}{=} \phi(Y_s)$, where ϕ is as guaranteed by the lemma. Hence, $X_s \subseteq \sigma_1(Y_s)$, and $|X_s| = |Y_s|$. Note that while all elements of Y_s belong to L_s , the elements of X_s are contained in several L_j 's, $j < s$.

For each j , $0 \leq j < s$, let $X_{s,j} \stackrel{\text{def}}{=} X_s \cap L_j$. Let $X_{s,r}$ be the largest such set. Since $|X_s| = |Y_s| \geq \frac{\epsilon_M(f)}{2^n} \cdot 2^n$, we have $|X_{s,r}| \geq \frac{\epsilon_M(f)}{2^{n^2}} \cdot 2^n$. Finally, let $Y_{s,r} \stackrel{\text{def}}{=} \phi^{-1}(X_{s,r})$. Then Lemma 7 holds with $S = Y_{s,r} \subseteq L_s$ and $R = X_{s,r} \subseteq L_r$. \blacksquare

Comment: To obtain the stronger bound on the sizes of S and R we do the following. Let $\text{dev} \stackrel{\text{def}}{=} \sqrt{\frac{1}{2}n \cdot \ln(8/\epsilon_M(f))}$. Then we have that the total number of strings in layers L_i where $i > \frac{n}{2} + \text{dev}$ is at most $\frac{\epsilon_M(f)}{8} \cdot 2^n$. Similarly, the total number of strings in layers L_i where $i < \frac{n}{2} - \text{dev}$ is at most $\frac{\epsilon_M(f)}{8} \cdot 2^n$. Assuming (without loss of generality) that $|D_1| \geq \frac{\epsilon_M(f)}{2} \cdot 2^n$, we have that the number of strings in D_1 that belong to layers L_i where $i \leq \frac{n}{2} + \text{dev}$ is at least $\frac{3\epsilon_M(f)}{8} \cdot 2^n$. By invoking Lemma 9 on the set $Y = D_1 \cap \left(\bigcup_{i \leq \frac{n}{2} + \text{dev}} L_i\right)$, we have a one-to-one mapping ϕ from Y to $X = \phi(Y) \subseteq \sigma_1(Y)$. Note that by definition of Y , $X \subseteq \bigcup_{i < \frac{n}{2} + \text{dev}} L_i$.

Since $|X| = |Y| \geq \frac{3\epsilon_M(f)}{8} \cdot 2^n$, and the total number of strings in layers L_i where $i < \frac{n}{2} - \text{dev}$ is at most $\frac{\epsilon_M(f)}{8} \cdot 2^n$, we have that $\left|X \cap \left(\bigcup_{i=\frac{n}{2}-\text{dev}}^{\frac{n}{2}+\text{dev}} L_i\right)\right| \geq \frac{\epsilon_M(f)}{4} \cdot 2^n$. For each i , $\frac{n}{2} - \text{dev} \leq i < \frac{n}{2} + \text{dev}$, let $X_i \stackrel{\text{def}}{=} X \cap L_i$ and let X_r be the largest such set. Then $|X_r| \geq \frac{\epsilon_M(f)}{4 \cdot 2^{\text{dev}}} \cdot 2^n$. Let $Y_r \stackrel{\text{def}}{=} \phi^{-1}(X_r)$, and for each i , $\frac{n}{2} - \text{dev} < i \leq \frac{n}{2} + \text{dev}$, define $Y_{i,r} \stackrel{\text{def}}{=} Y_r \cap L_i$. Then there exists a set $Y_{s,r} \subseteq L_s$ such that $|Y_{s,r}| \geq \frac{\epsilon_M(f)}{16 \cdot \text{dev}^2} \cdot 2^n = \Omega\left(\frac{\epsilon_M(f)}{n \cdot \log(1/\epsilon_M(f))}\right) \cdot 2^n$. We then let $S = Y_{s,r}$ and $R = X_{s,r}$.

3.2 Existence of disjoint paths between matched sets

Let $S \subseteq L_s$ and $R \subseteq L_r$ be as stated in Lemma 8, and let $d = s - r$. Recall that for each $0 \leq i \leq n$, L_i is the set of all vertices in G_n corresponding to strings with exactly i 1's. We shall prove Lemma 8 by induction on m and d . The base cases, i.e., the case where $m = 1$ and $d \geq 1$, and the case where $d = 1$ and $m \geq 1$, clearly hold. Consider general $m > 1$ and $d > 1$, and assume by induction that the claim holds for every pair m' and d' such that either $m' < m$ and $d' \leq d$ or $m' \leq m$ and $d' < d$. Let Q be the set of vertices in L_{s-1} that are on a directed path going from some vertex in S to some vertex in R , and let P be the set of vertices in L_{r+1} that are on such directed paths from S to R (see Figure 3.2). We shall prove the induction claim in two steps. In the first step we use the induction hypothesis (for $m' < m$ and $d' = d$) to show that either $|Q| \geq m$ or $|P| \geq m$ (or both). In the second step we use this fact together with the induction hypothesis (for $m' < m$ and $d' = d$ and for $m' = m$ and $d' < d$) to prove the induction claim.

Step 1: Either $|Q| \geq m$ or $|P| \geq m$.

Proof: Consider the subgraph G'_n of G_n containing S , R and all vertices and edges that belong to paths between S and R .

Claim 8.1: *Let v be a vertex in S and let u be a vertex in some level L_i , where $r+1 \leq i \leq s-1$, such that there is a directed path from v to u in G'_n . Then the outdegree of v in G'_n is at*




Figure 2. The sets S, R, Q, and P.

least as large as the outdegree of u in G'_n . Similarly, if $w \in R$ and $z \in L_i$ where $r + 1 \leq i \leq s - 1$, such that there is a directed path from z to w in G'_n , then the indegree of w in G'_n is at least as large as the indegree of z in G'_n .

Proof: We prove the claim concerning outdegrees. The claim about indegrees is proved analogously. Let $a \geq 1$ be the outdegree of u and consider the vertices u^1, \dots, u^a in L_{i-1} such that there is an edge in G'_n from u to each u^i . Recall that by definition of G'_n there are paths in G'_n from the u^i 's to vertices in R . Therefore, any vertex that is on a path from v to one of the u^i 's is in G'_n as well.

For each u^i , let $b^i \in [n]$ be the index of the bit on which the strings corresponding to u and u^i differ; i.e., $u_{b^i} = 1$ while $u^i_{b^i} = 0$. By definition, b^1, \dots, b^a are distinct indices, and since $u_{b^i} = 1$ for every i , it also holds that $v_{b^i} = 1$ for every i . For each b^i , let v^i be the vertex in L_{s-1} that differs from v on the b^i 'th bit; i.e., $v^i_{b^i} = 0$, and for every $j \neq b^i$, $v^i_j = v_j$. Then each of the a v^i 's is on a path from v to u^i , and the claim follows. \square

We note that the above claim can be strengthened to show that the outdegree of v (respectively, indegree of w) is greater than the outdegree of u (respectively, indegree of z), by at least $s - i$ (respectively, $i - r$). This is done by taking into account the bits on which v and u (respectively, w and z) differ.

Let k be the maximum outdegree of vertices in S , and let t be the maximum indegree of vertices in R . We partition S , R , Q , and P into subsets according to their degrees in G'_n as follows. For every $i \leq k$ we let S_i be the subset of vertices in S that have outdegree exactly i , and for every $j \leq t$, we let R^j be the subset of vertices in R that have indegree exactly j . Similarly we let Q_i^j (respectively, P_i^j) be the subset of vertices in Q (respectively, P) with outdegree exactly i and indegree exactly j . First note that by Claim 8.1, the maximum outdegree of vertices in Q and P is at most k , and the maximum indegree is at most t . Therefore, for every j and $i > k$, $|Q_i^j|, |P_i^j| = 0$,

and for every i and $j > t$, $|Q_i^j|, |P_i^j| = 0$.

Furthermore, by Claim 8.1, for every i , and each vertex $v \in S_i$, the vertices u in P such that there exists a directed path from v to u must belong to $\cup_{i' \leq i} \cup_{j' \leq j} P_{i', j'}$. For any $q \leq k$, let $S_{\leq q} = \cup_{i=1}^q S_i$. By definition of k (as the maximum degree of vertices in S), the set S_k is nonempty and hence for every $q < k$, $|S_{\leq q}| < m$. Therefore, we can apply the induction hypothesis and obtain that there exist vertex disjoint paths between $S_{\leq q}$ and $\phi(S_{\leq q})$ (where ϕ is the matching guaranteed by the hypothesis of Lemma 8). For any $q < k$ let $\Pi(S_{\leq q}) \subseteq Q$ denote the set of neighbors of vertices in $S_{\leq q}$ that lie on these paths to $\phi(S_{\leq q})$. Since these paths are disjoint, $|\Pi(S_{\leq q})| = |S_{\leq q}|$. Using the above and the fact that the S_i 's are disjoint and the P_i^j 's are disjoint, the following inequality holds for every $q < k$:

$$\begin{aligned} \sum_{i=1}^q |S_i| &= |S_{\leq q}| = |\Pi(S_{\leq q})| \leq \left| \cup_{i=1}^q \cup_{j=1}^{\infty} P_i^j \right| \\ &= \sum_{i=1}^q \sum_{j=1}^{\infty} |P_i^j| = \sum_{i=1}^q \sum_{j=1}^k |P_i^j|. \end{aligned} \quad (5)$$

Similarly, we can obtain that for every $p < s$,

$$\sum_{j=1}^p |R^j| \leq \sum_{i=1}^k \sum_{j=1}^p |Q_i^j|. \quad (6)$$

Recall that we would like to show that either $|Q| \geq m$ or $|P| \geq m$. Thus, assume in contradiction that both $|Q| < m$ and $|P| < m$. Therefore, by Equation (5), for every $q < k$,

$$\begin{aligned} \sum_{i=q+1}^k |S_i| &= |S| - \sum_{i=1}^q |S_i| = m - \sum_{i=1}^q |S_i| \\ &\geq m - \sum_{i=1}^q \sum_{j=1}^t |P_i^j| \\ &> |P| - \sum_{i=1}^q \sum_{j=1}^t |P_i^j| \end{aligned} \quad (7)$$

and so

$$\sum_{i=q+1}^k |S_i| > \sum_{i=q+1}^k \sum_{j=1}^t |P_i^j|. \quad (8)$$

Similarly, for every $p < t$,

$$\sum_{j=p+1}^t |R^j| > \sum_{i=1}^k \sum_{j=p+1}^t |Q_i^j|. \quad (9)$$

By summing both sides of Equation (8) over all $q < k$ we get

$$\sum_{q=0}^{k-1} \sum_{i=q+1}^k |S_i| > \sum_{q=0}^{k-1} \sum_{i=q+1}^k \sum_{j=1}^t |P_i^j| \quad (10)$$

or equivalently,

$$\sum_{i=1}^k i \cdot |S_i| > \sum_{i=1}^k \sum_{j=1}^t i \cdot |P_i^j|. \quad (11)$$

Similarly, from Equation (9) we get

$$\sum_{j=1}^t j \cdot |R^j| > \sum_{i=1}^k \sum_{j=1}^t j \cdot |Q_i^j|. \quad (12)$$

Summing Equations (11) and (12), we get

$$\sum_{i=1}^k i \cdot |S_i| + \sum_{j=1}^t j \cdot |R^j| > \sum_{i=1}^k \sum_{j=1}^t i \cdot |P_i^j| + \sum_{i=1}^k \sum_{j=1}^t j \cdot |Q_i^j|. \quad (13)$$

However, since the number of edges going out of vertices in S equals the number of edges entering vertices in Q we have that:

$$\sum_{i=1}^k i \cdot |S_i| = \sum_{i=1}^k \sum_{j=1}^t j \cdot |Q_i^j| \quad (14)$$

and similarly for R and P we have

$$\sum_{j=1}^t j \cdot |R^j| = \sum_{i=1}^k \sum_{j=1}^t i \cdot |P_i^j|. \quad (15)$$

Summing Equations (14) and (15) we get

$$\sum_{i=1}^k i \cdot |S_i| + \sum_{j=1}^t j \cdot |R^j| = \sum_{i=1}^k \sum_{j=1}^t j \cdot |Q_i^j| + \sum_{i=1}^k \sum_{j=1}^t i \cdot |P_i^j| \quad (16)$$

contradicting Equation (13). ■ (Step 1.)

Step 2: There exist vertex disjoint paths from S to R .

Proof: From Step 1 we have that either $|Q| \geq m$ or $|P| \geq m$. Assume the former is true — we shall see that this can be done without loss of generality. We next show that (1) there exists a perfect matching between S and (a subset of) Q ; and (2) there exists a 1-to-1 mapping ϕ' from the matched vertices of Q to R so that there is a path from each matched $u \in Q$ to $\phi'(u)$. Given (2) we can apply the induction hypothesis for $d' = d - 1$ (and $m' = m$) on Q and R , and by combining with (1) we get the desired paths from S to R .

We actually prove both (1) and (2) together. Consider the following auxiliary network, A . It has a single source vertex s , a single target vertex t , and the rest of the vertices are partitioned into three layers corresponding to S , Q and R , respectively. There is an edge from s to each of the vertices in S , and from each of the vertices in R to t . The edges between S and Q are as in G'_n and edges between Q and R correspond to directed paths in G'_n . We show that the minimum $s - t$ vertex-separator in A has size m . Items (1) and (2) follow by

one of the variations of Menger's Theorem (see [15, Thm. 6.4 and discussion on pp. 130]), which guarantees the existence of m vertex-disjoint paths from s to t .

Assume in contradiction that there exists a vertex-separator C of size smaller than m in A . Let $m_1 \stackrel{\text{def}}{=} |C \cap S|$, $m_2 \stackrel{\text{def}}{=} |C \cap Q|$, and $m_3 \stackrel{\text{def}}{=} |C \cap R|$. Consider the subset of vertices $S' \subseteq S$ that do not belong to C and are not mapped by ϕ to vertices in $R \cap C$. The size of S' is at least $m' = m - (m_1 + m_3) > |C| - (m_1 + m_3) = m_2$. Let $R' \stackrel{\text{def}}{=} \phi(S')$, and Q' be the subset of vertices in Q that are on a directed path in G'_n going from some vertex in S' to a vertex in R' .

We consider two cases. If $S' = S$ (i.e., $C \subseteq Q$) then $Q' = Q$, and since $|C| < m \leq |Q|$, there exists at least one vertex in $Q \setminus C$ on a path from a vertex in S to a vertex in R , contradicting the assumption that C is a vertex separator. If $S' \subset S$, then by the induction hypothesis (for $m' = |S'| < m$ and $d' = d$), there exist vertex disjoint paths in G'_n from S' to $\phi(S')$ and hence necessarily $|Q'| \geq |S'| > m_2$. Since $|C \cap Q| = m_2$, we again reach contradiction to the assumption that C is a vertex separator. ■

4 Testing whether a function is unate

By our definition of monotonicity (used throughout the paper), a function is said to be monotone if, for any string, flipping any bit of the string from 0 to 1, does not decrease the value of the function. A more general notion is that of *unate* functions. A function f is *unate* if there exists a string $\zeta = \zeta_1 \dots \zeta_n \in \{0, 1\}^n$ for which the following holds: For any string $x = x_1 \dots x_n$, and for any i such that $x_i = \zeta_i$, if we let $y = x_1, \dots, x_{i-1}, \neg x_i, x_{i+1}, \dots, x_n$ (i.e., y is the same as x except for the i^{th} bit, which is flipped from ζ_i to $\neg \zeta_i$), then $f(y) \geq f(x)$. We say in such a case the f is *monotone with respect to* ζ . In particular, if a function is monotone with respect to the all-0 string, then we simply say that it is a monotone function, and if a function is monotone with respect to *some* ζ , then it is unate. Thus, the generalization of monotonicity to unateness allows that for each position there be a (possibly different) *direction* (i.e., not necessarily the $0 \rightarrow 1$ direction), such that the value of the function cannot decrease when the bit is flipped in that direction.

Similarly to Algorithm 1 (for testing monotonicity), which searches for evidence to non-monotonicity, the testing algorithm for unateness tries to find evidence to non-unateness. However, here it does not suffice to find a pair of strings x, y that differ on the i^{th} bit such that $x < y$ while $f(x) > f(y)$, since f could be monotone with respect to ζ such that $\zeta_i = 1$. Instead we search for *two pairs* of strings, $x^1 < y^1$ and $x^2 < y^2$, where each pair differs on the (same) i^{th} bit, such that $f(x^1) > f(y^1)$ and $f(x^2) < f(y^2)$ (or vica versa). This implies that there is no ζ such that f is monotone with respect to ζ (since, in particular, ζ_i can be neither 0 nor 1).

ALGORITHM 3 (TESTING UNATENESS): On input n, ϵ and oracle access to $f : \{0, 1\}^n \mapsto \{0, 1\}$, do the following:

1. Uniformly select $m = O(n^{3.5}/\epsilon)$ strings in $\{0, 1\}^n$, denoted x^1, \dots, x^m , and m indices in $\{1, \dots, n\}$, denoted i^1, \dots, i^m .
2. For each selected x^j , obtain the values of $f(x^j)$ and $f(y^j)$, where y^j results from x^j by flipping the i^j 'th bit.
3. If unateness is found to be violated then reject.

Violation occurs, if among the string-pairs $\{x^j, y^j\}$, there exist two pairs and an index i , such that in both pairs the strings differ on the i^{th} bit, but in one pair the value of the function increases when the bit is flipped for 0 to 1, and in the other pair the value of the function increases when the bit is flipped from 1 to 0.

If no contradiction to unateness was found then accept.

Theorem 10 *Algorithm 3 is a testing algorithm for unateness. Furthermore, if the function is unate, then Algorithm 3 always accepts.*

We shall need the following notation. For $\zeta \in \{0, 1\}^n$, let \prec_ζ denote the partial order on strings with respect to ζ . Namely, $x \prec_\zeta y$ if and only if $x \oplus \zeta \prec y \oplus \zeta$. Let $\epsilon_{M,\zeta}(f)$ denote the minimum distance between f and any function g that is monotone with respect to ζ , and let $\delta_{M,\zeta}(f)$ denote the fraction of pairs x, y that differ on a single bit such that $x \prec_\zeta y$ but $f(x) > f(y)$. It follows from the above definitions that for any f and ζ , $\epsilon_{M,\zeta}(f) = \epsilon_M(f_\zeta)$ and $\delta_{M,\zeta}(f) = \delta_M(f_\zeta)$, where f_ζ is defined by $f_\zeta(x) = f(x \oplus \zeta)$. Hence, as a corollary to Theorem 2, we have

Corollary 11 *For any $f : \{0, 1\}^n \mapsto \{0, 1\}$, and for any $\zeta \in \{0, 1\}^n$, $\delta_{M,\zeta}(f) \geq \frac{\epsilon_{M,\zeta}(f)}{n^3}$.*

Proof of Theorem 10: For each $i \in \{1, \dots, n\}$, let $\gamma_{i,0}(f)$ denote the fraction, among all pairs of strings that differ on a single bit, of the pairs x, y such that x and y differ only on the i^{th} bit, $x_i = 0$, $y_i = 1$, and $f(x) > f(y)$. Similarly, let $\gamma_{i,1}(f)$ denote the fraction of pairs of strings x, y such that x and y differ only on the i^{th} bit, $x_i = 1$, $y_i = 0$, and $f(x) > f(y)$. In other words, $\gamma_{i,0}(f)$ is the fraction of pairs that can serve as evidence to f not being monotone with respect to any ζ such that $\zeta_i = 0$, while $\gamma_{i,1}(f)$ is the fraction of pairs that can serve as evidence to f not being monotone with respect to any ζ such that $\zeta_i = 1$. Note that in case f is monotone with respect to some ζ , then for every i , $\gamma_{i,\zeta_i}(f) = 0$. More generally, $\delta_{M,\zeta}(f) = \sum_{i=1}^n \gamma_{i,\zeta_i}(f)$ holds for every $\zeta \in \{0, 1\}^n$ (since each edge contributing to $\delta_{M,\zeta}(f)$ contributes to exactly one γ_{i,ζ_i}).

Let us define $\epsilon_U(f)$ to be $\min_\zeta(\epsilon_{M,\zeta}(f))$ so that it equals the minimum distance of f to any unate function (i.e., any function that is monotone with respect to some ζ).

Claim 10.1. $\sum_{i=1}^n \min(\gamma_{i,0}(f), \gamma_{i,1}(f)) \geq \frac{\epsilon_U(f)}{n^3}$.

Proof: Let $\zeta = \zeta_1 \dots \zeta_n$ be defined as follows: For each i , if $\gamma_{i,0}(f) \leq \gamma_{i,1}(f)$ then $\zeta_i = 0$, and otherwise, $\zeta_i = 1$. In other words, $\zeta_i = \arg\min_{b \in \{0,1\}}(\gamma_{i,b})$. The key observation is $\delta_{M,\zeta}(f) = \sum_{i=1}^n \gamma_{i,\zeta_i} = \sum_{i=1}^n \min(\gamma_{i,0}(f), \gamma_{i,1}(f))$, where the first equality holds for any ζ , and the second follows from the definition of this specific ζ . Invoking Corollary 11, we have $\delta_{M,\zeta}(f) \geq \frac{\epsilon_{M,\zeta}(f)}{n^3} \geq \frac{\epsilon_U(f)}{n^3}$. \square

For each i , let $\Gamma_{i,0}(f)$ be the set of all pairs of strings x, y that differ only on the i^{th} bit, where $x_i = 0$ and $y_i = 1$, and such that $f(x) > f(y)$. Similarly, let $\Gamma_{i,1}(f)$ be the set of all pairs x, y that differ only on the i^{th} bit, where $x_i = 1$ and $y_i = 0$, and such that $f(x) > f(y)$. Claim 10.1 gives us a lower bound on the sum $\sum_i \min(|\Gamma_{i,0}|, |\Gamma_{i,1}|)$. To prove Theorem 10, it suffices to show that if we uniformly select $\Omega(n^{3.5}/\epsilon_U(f))$ pairs of strings that differ on a single bit, then with probability at least $2/3$, for some i we shall obtain both a pair belonging to $\Gamma_{i,0}(f)$ and a pair belonging to $\Gamma_{i,1}(f)$. The above is derived from the following technical claim, which can be viewed as a generalization of the *Birthday Paradox*, and whose proof is given in the full version of this paper [18].

Claim 10.2. *Let $S_1, \dots, S_n, T_1, \dots, T_n$ be disjoint sets of elements belonging to domain X . For each i , let the probability of selecting an element x in S_i (when x is chosen uniformly in X), be p_i , and the probability of selecting an element in T_i , be q_i . Suppose that for all i , $q_i \geq p_i$, and that $\sum_i p_i \geq \rho$ for some $\rho > 0$. Then, for some constant c , if we uniformly select $c \cdot \sqrt{n}/\rho$ elements in X , then with probability at least $2/3$, for some i we shall obtain one element in S_i and one in T_i .*

Acknowledgments

We would like to thank Dan Kleitmann for a very helpful discussion.

References

- [1] D. Angluin. Queries and concept learning. *Machine Learning*, 2(4):319–342, April 1988.
- [2] S. Arora, C. Lund, R. Motwani, M. Sudan, and M. Szegedy. Proof verification and intractability of approximation problems. *JACM*, 45(3):501–555, 1998.
- [3] S. Arora and S. Safra. Probabilistic checkable proofs: A new characterization of NP. *JACM*, 45(1):70–122, 1998.
- [4] S. Arora and S. Sudan. Improved low degree testing and its applications. In *Proceedings of STOC97*, pages 485–495, 1997.

- [5] L. Babai, L. Fortnow, L. Levin, and M. Szegedy. Checking computations in polylogarithmic time. In *Proceedings of STOC91*, pages 21–31, 1991.
- [6] L. Babai, L. Fortnow, and C. Lund. Non-deterministic exponential time has two-prover interactive protocols. *Computational Complexity*, 1(1):3–40, 1991.
- [7] M. Bellare, D. Coppersmith, J. Håstad, M. Kiwi, and M. Sudan. Linearity testing in characteristic two. In *Proceedings of FOCS95*, pages 432–441, 1995.
- [8] M. Bellare, O. Goldreich, and M. Sudan. Free bits, PCPs and non-approximability – towards tight results. *SIAM Journal on Computing*, 27(3):804–915, 1998.
- [9] M. Bellare, S. Goldwasser, C. Lund, and A. Russell. Efficient probabilistically checkable proofs and applications to approximation. In *Proceedings of STOC93*, pages 294–304, 1993.
- [10] M. Bellare and M. Sudan. Improved non-approximability results. In *Proceedings of STOC94*, pages 184–193, 1994.
- [11] A. Blum, C. Burch, and J. Langford. On learning monotone Boolean functions. In *Proceedings of FOCS98*, 1998.
- [12] M. Blum, M. Luby, and R. Rubinfeld. Self-testing/correcting with applications to numerical problems. *JACM*, 47:549–595, 1993.
- [13] B. Bollobás. *Combinatorics*. Cambridge University Press, 1986.
- [14] F. Ergun, S. Kannan, S. R. Kumar, R. Rubinfeld, and M. Viswanathan. Spot-checkers. In *Proceedings of STOC98*, pages 259–268, 1998.
- [15] S. Even. *Graph Algorithms*. Computer Science Press, 1979.
- [16] U. Feige, S. Goldwasser, L. Lovász, S. Safra, and M. Szegedy. Approximating clique is almost NP-complete. *JACM*, 43(2):268–292, 1996.
- [17] P. Gemmell, R. Lipton, R. Rubinfeld, M. Sudan, and A. Wigderson. Self-testing/correcting for polynomials and for approximate functions. In *Proceedings of STOC91*, pages 32–42, 1991.
- [18] O. Goldreich, S. Goldwasser, E. Lehman, and D. Ron. Testing monotonicity. Available from <http://theory.lcs.mit.edu/~danar>, 1998.
- [19] O. Goldreich, S. Goldwasser, and D. Ron. Property testing and its connection to learning and approximation. In *Proceedings of FOCS96*, pages 339–348, 1996.
- [20] O. Goldreich and D. Ron. Property testing in bounded degree graphs. In *Proceedings of STOC97*, pages 406–415, 1997.
- [21] O. Goldreich and D. Ron. A sublinear bipartite tester for bounded degree graphs. In *Proceedings of STOC98*, 1998.
- [22] J. Håstad. Testing of the long code and hardness for clique. In *Proceedings of STOC96*, pages 11–19, 1996.
- [23] J. Håstad. Getting optimal in-approximability results. In *Proceedings of STOC97*, pages 1–10, 1997.
- [24] M. Kearns, M. Li, and L. Valiant. Learning boolean formulae. *JACM*, 41(6):1298–1328, 1994.
- [25] M. Kearns and L. Valiant. Cryptographic limitations on learning boolean formulae and finite automata. *JACM*, 41(1):67–95, 1994.
- [26] M. Kiwi. *Probabilistically Checkable Proofs and the Testing of Hadamard-like Codes*. PhD thesis, MIT, 1996.
- [27] R. Raz and S. Safra. A sub-constant error-probability low-degree test, and a sub-constant error-probability PCP characterization of NP. In *Proceedings of STOC97*, pages 475–484, 1997.
- [28] R. Rubinfeld and M. Sudan. Robust characterization of polynomials with applications to program testing. *SIAM Journal on Computing*, 25(2):252–271, 1996.
- [29] L. Trevisan. Recycling queries in pcps and in linearity tests. In *Proceedings of STOC98*, pages 299–308, 1998.